

CS 460/560

Introduction to Computational Robotics
Fall 2019, Rutgers University

Lecture 18

Feedback & Potential Field Based Planners

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Instructor: Jingjin Yu

Outline

Artificial potential field


Feedback based planner

- ⇒ Discrete state spaces

- ⇒ Vector fields for continuous domains

Potential Field

Potential fields in the nature



⇒ A dropped ball or free water will simply fall with gravity

⇒ Same applies to other types of potential fields

⇒ E.g., magnets can be used to pick up small ferrous metals (iron, steel)

⇒ E.g., the forces between gaseous molecules

⇒ There can be attractive and repulsive forces

⇒ Recall in the case of gravity, the potential $U(r) = -\frac{GMm}{r}$


⇒ M, m are the masses of the two attracting bodies, G is a constant, r is the distance

⇒ $F = -\frac{dU}{dr} = -\frac{GMm}{r^2}$ ($\frac{GM}{r^2}$ is about 9.8 m/s^2 near earth's surface)

Artificial Potential Field for Motion Planning (I)

Potential fields inspired the development of **artificial potential field** for motion planning. General idea:


- ⇒ Used to guide a robot to a goal region from anywhere in the environment
- ⇒ To do so, builds a **navigation function** for the robot to follow
- ⇒ The function can be thought of forces from an artificial potential field
- ⇒ Includes an **attraction term** to pull the robot toward the goal
- ⇒ Includes **repulsive terms** to push the robot away from obstacles
- ⇒ Together, creating a desired potential field
- ⇒ This yields a **function** $U(x)$



Artificial Potential Field for Motion Planning (II)

Given a potential field $U(x)$, to reach goal from configuration x


- ⇒ Compute the negative gradient $-\nabla U = -\left[\frac{\partial U}{\partial x_1}, \dots, \frac{\partial U}{\partial x_n}\right]$
- ⇒ Simply follow it! Update as x changes
- ⇒ $U(x)$ may be viewed as level sets



Constructing an artificial potential field can be tricky

- ⇒ Generally easier to do for convex obstacles
- ⇒ Can be challenging for non-convex obstacles
 - ⇒ May create local minima
 - ⇒ Traps the robot
- ⇒ There are ways to avoid it for “star worlds”
- ⇒ Applies to higher dimensions as well
- ⇒ For details, see Rimon & Kodischeck, '92

May be viewed a **feedback** based planner



Problems with Classic Path Planners

Most planners covered until now are **open loop**

- ⇒ Paths generated by these planners are for ideal conditions
 - ⇒ E.g., applying same thrust to two wheels of a DDR, the robot will move straight
 - ⇒ Works fine for ideal models, e.g., simulation, games
 - ⇒ What happens if you do this to the our robot?
 - ⇒ Most likely not a straight line
 - ⇒ Otherwise, our lives will be much simpler!
 - ⇒ Perhaps boring though...
- ⇒ How to accommodate such variations?
- ⇒ Use **feedback!**
 - ⇒ Potential field is such a planner
- ⇒ Feedback allows the handling of many sources of uncertainties, e.g.
 - ⇒ Motor speed variations between different motors under the same input
 - ⇒ Motor speed changes as battery charging state changes
 - ⇒ Uneven terrain: carpet, wooden floor, gravel, ...
 - ⇒ Air flow
 - ⇒ ...



Feedback Based Planner



What does feedback based planner do?

⇒ Assumptions

- ⇒ The robot has a good estimate of its current state $x \in X$
- ⇒ There is a specific goal set (may be multiple goals) to navigate to

⇒ Feedback plan is actually a **policy** $\pi(x)$

- ⇒ At each $x \in X$, $\pi: X \rightarrow U, x \mapsto \pi(x)$ produces a control for the robot
- ⇒ All states in the free space must be covered
- ⇒ The policy is executed until a goal is reached




- ⇒ Key: feedback based planner provides a (lookup) policy. In particular, **no search** is needed after the policy is created.

Finding Feasible Solutions

Feedback based problem uses transition function $f(x, u)$

⇒ For a finite control set U , this defines a directed graph. How?

⇒ Basically, $x' = f(x, u)$ induces a directed edge





⇒ To find a policy, we need to grow a directed acyclic graph (DAG) from x_G

⇒ To do this, first reverse all edge directions

⇒ Then execute any **spanning tree** algorithm

⇒ This yields a spanning tree from x_G to all other states in the free space



⇒ Reversing the edge directions again then yields a feasible solution



Finding Optimal Solutions

What about optimal feedback policy?

- ⇒ We may simply run Dijkstra's algorithm!
 - ⇒ There are other faster algorithms, but with similar principle
- ⇒ Since we inverted the edges, how do we know this works?
 - ⇒ Can prove via contradiction
- ⇒ We can go one step further and build **navigation functions**
 - ⇒ Associate each state with a (cost-to-go) value
 - ⇒ For a non-goal node, there must exist a neighbor with lower value
 - ⇒ Essentially, this yields a level set (potential function)
 - ⇒ Can also take other forms, e.g., maximum clearance



22	21	22	21	20	19	18	17	16	17
21	20							15	16
20	19							14	15
19	18	17	16	15	14	13	12	13	14
18	17	16	15	14	13	12	11	12	13
								10	11
								9	10
3	2	1	2	3				8	9
2	1	0	1	2				7	8
3	2	1	2	3	4	5	6	7	8

Continuous Domain: Vector Field

Continuous domain is similar to the discrete domain but

⇒ Need to express the (control) policy differently

⇒ For this we need vector fields


Basic properties of vector fields are:

1. (Commutative Group Under Vector Addition) The set V is a commutative group with respect to vector addition, $+$.
2. (Associativity of Scalar Multiplication) For any $v \in V$ and any $\alpha, \beta \in \mathbb{F}$, $\alpha(\beta v) = (\alpha\beta)v$.
3. (Distributivity of Scalar Sums) For any $v \in V$ and any $\alpha, \beta \in \mathbb{F}$, $(\alpha + \beta)v = \alpha v + \beta v$.
4. (Distributivity of Vector Sums) For any $v, w \in V$ and any $\alpha \in \mathbb{F}$, $\alpha(v + w) = \alpha v + \alpha w$.
5. (Scalar Multiplication Identity) For any $v \in V$, $1v = v$ for the multiplicative identity $1 \in \mathbb{F}$.

How does it work?

⇒ Following smooth vector fields creates a smooth trajectory (integral curve)



⇒ If the vector field “flows” into the goal region, then we are done



Vector Field Construction Methods (I)

Triangulation based construction







- ⇒ First, triangulate the domain
- ⇒ Build a graph of the triangles and then build a spanning tree from x_G
- ⇒ For each triangle, generate a simple vector field pointing to the **exit edge**
- ⇒ Works for high dimensions



Vector Field Construction Methods (II)

Visibility based construction

- ⇒ Compute visibility region of goal
- ⇒ Do wavefront propagation
- ⇒ Iteratively applies to way points
- ⇒ The method is optimal
- ⇒ But difficult for high dimensions



Vector Field Construction Methods (III)

Sampling-based funnel composition

- ⇒ Sampling based methods also apply here!
- ⇒ First, compute a cover via sampling
- ⇒ Then, again build a graph based on cover adjacency
 - ⇒ I.e., two intersecting covers will share an edge
- ⇒ Obtain a spanning tree from the graph
- ⇒ For each pair of covers, construct a local vector field
- ⇒ All together, these become a bunch of “funnels”

