



CS 460/560

Introduction to Computational Robotics
Fall 2019, Rutgers University

Lecture 14

Intro To Sampling-Based Planning Methods (2)

| | | | |
|----|----|----|----|
| 1 | 2 | 3 | 4 |
| 5 | 6 | 7 | 8 |
| 9 | 10 | 11 | 12 |
| 13 | 14 | 15 | |

Instructor: Jingjin Yu

Outline (for Next 3-4 Lectures)

Drawbacks of combinatorial motion planning methods

Probabilistic roadmap (PRM) introduction

Components of sampling-based motion planning methods

- ⇒ Sampling
- ⇒ k -d tree and nearest neighbor search
- ⇒ Distance metric
- ⇒ Collision detection

PRM in more detail

A new notion of completeness

Rapidly-exploring random trees (RRT)

When would sampling-based method work well?


Optimality issues

Drawbacks of Combinatorial Methods


Recall the illustration of the Piano Mover's Problem

- ⇒ Modeling of the free configuration space C_{free} can be a daunting task – they have to be represented as semi-algebraic sets
- ⇒ The associated computation is also prohibitive for even just a few degrees of freedom

To the rescue: **sampling based methods** – instead of representing C_{free} explicitly and globally, we instead “probe” the space locally, as necessary




Sampling-Based Planning



Key Components of Sampling-Based Planning

Sampling-based planning requires several important subroutines


- ⇒ An **efficient sampling routine** is needed to generate the samples. These samples should **cover** C_{free} well in order to be effective
- ⇒ **Efficient nearest neighbor search** is necessary for quickly building the roadmap: for each sample in C_{free} we must find its k -nearest neighbors
- ⇒ The neighbor search also requires a **distance metric** to be properly defined so we know the distance between two samples
 - ⇒ This can be tricky for certain spaces, e.g., $SE(3)$
- ⇒ **Collision checking** - Note that C_{free} is not computed explicitly so we actually are checking collisions between a complex robot and a complex environment



Nearest Neighbor Search


Connecting the samples

- ⇒ Building the graph requires connecting the samples
- ⇒ This cannot be done for all pairs of points!
 - ⇒ For N sample points, this requires N^2 operations
 - ⇒ But N can be very large, e.g., $> 10^5$
 - ⇒ For $N = 10^6$, $N^2 = 10^{12}$
- ⇒ We have to do it more efficiently!
- ⇒ This is known as nearest neighbor search



Variants of useful nearest neighbor search


- ⇒ 1-NN: finding a single nearest neighbor
 - ⇒ Can be done with k -d trees
 - ⇒ We will look at this in more detail
- ⇒ k -NN: finding k nearest neighbors
 - ⇒ Note the k here is not the same as the k in k -d trees
 - ⇒ Can run 1-NN algorithms k times



k -d Tree


k -d tree stands for **k -dimensional trees**

- ⇒ A data structure for storing points in k dimensions
- ⇒ Assumes a tree like structure
- ⇒ Useful for finding points w/ certain properties
- ⇒ Can be used for solving 1-NN



Construction of a k -d tree for n points

- ⇒ Pick dimension i , pick a point with coordinates $x = (x_1, \dots, x_i, \dots, x_k)$
- ⇒ Split the points based on x_i (greater or less than)
- ⇒ Repeat the above two steps recursively
 - ⇒ Increase i (modulo k) each time
 - ⇒ I.e., pick a new dimension each time
- ⇒ Depth: $\log n$ if balanced
- ⇒ Construction takes $O(kn \log n)$ time
 - ⇒ Each dimension needs sorting $\sim O(n \log n)$
- ⇒ Can speed up to $O(n \log n)$
- ⇒ Balancing is important






Nearest Neighbor Search w/ k -d Tree

Finding nearest neighbor of a query point q

- ⇒ Basically, traverse the tree, e.g., using BFS
- ⇒ Maintain a current best **candidate** x
- ⇒ Also maintain a queue of subtree distances to q
- ⇒ Uses the subtree distances to prioritize search

Example

- ⇒ Start with root $(0.42, 0.42)$
 - ⇒ $x = (0.42, 0.42)$, both left and right subtrees are active
- ⇒ Examine $(0.07, 0.55)$
 - ⇒ Three trees on the queue afterward
- ⇒ Examine $(0.82, 0.3)$, update $x = (0.82, 0.3)$
 - ⇒ Truncate the left subtrees
 - ⇒ Two subtrees left
- ⇒ Examine $(0.72, 0.68)$, update $x = (0.72, 0.68)$
 - ⇒ We are done since the last subtree is further from q than x



Performance of k -d Tree and Generalization

General performance of balanced k -d tree


- ⇒ Construction: $O(n \log n)$ w/ $O(n)$ median computation
- ⇒ Construction through presorting the points: $O(kn \log n)$
- ⇒ Inserting/deletion of a new point: $O(\log n)$
- ⇒ Nearest neighbor search: $O(\log n)$ for randomly distributed points

k -d trees can be used for k -NN (k nearest neighbor search) as well

- ⇒ Naïve implementation: simply run 1-NN k times
- ⇒ This yields $O(k \log n)$ running time
- ⇒ Improvement
 - ⇒ Keep up to k candidates
 - ⇒ Only discard a subtree if worse than all k candidates

Applications of k -NN

- ⇒ Widely used in classification tasks, e.g.,
 - ⇒ Optical character recognition (OCR)
 - ⇒ Pattern recognition



A Brief Look at the Issue of Distance Metric

Nearest neighbor queries requires a distance metric

⇒ Given two points x and q , need to know their distance $d(x, q)$

⇒ Otherwise, cannot compare!

⇒ This is easy in Euclidean space: $d(x, q) = \|x - q\|_2 = \sqrt{\sum (x_i - q_i)^2}$

⇒ But what about T^2 or $\mathbb{R}^2 \times S^1$, or more complex settings?



⇒ For T^2 , θ_1 seems to be more important

⇒ For $\mathbb{R}^2 \times S^1$, a small change in θ can be hard to make

⇒ Sometimes, we can work with the workspace or task space

⇒ This however will make the sampling more difficult

⇒ There is no universal solution – requires some creativity




Collision Detection

Sampling based methods need to check whether robot is in collision


⇒ Generally, given two sets of points A and B , we want to check the distance between them

$$d(A, B) = \min_{a \in A, b \in B} |a - b|$$

⇒ Clearly, A and B intersect (collide) if and only if $d(A, B) = 0$




$d(A, B) > 0$: no collision



$d(A, B) = 0$: collision

Bounded Volume Hierarchy (BVH)

Collision checking can be difficult for general objects, e.g.,



$d(A, B)$ are hard to compute directly!


Often, simpler **bounding volumes** are used to approximate the shapes

- ⇒ However, bounding volumes **over approximate** the shapes
- ⇒ No collision between bounding volumes → no collision between the shapes
- ⇒ Collision between bounding volumes → **possible** collision
- ⇒ Need to refine hierarchically if a possible collision is detected
- ⇒ Such a method is called **bounded volume hierarchy** (BVH)

Bounded Volume Hierarchy (BVH)

BVH breaks objects into smaller pieces




Which yields a hierarchy, represented as a **tree**



This is carried out incrementally

⇒ Finer hierarchies are created as needed and then saved for later

Bounded Volume Hierarchy (BVH), Continued



For collision checking, it works with two BVH trees

- ⇒ Starting from the roots and check for collision (how?)
 - ⇒ No collision → done with the branch
 - ⇒ Otherwise, check pairs of children on the trees
- ⇒ Recursively call the procedure
- ⇒ Traverse down the tree
- ⇒ How many possible checks in total (say each object has n pieces)?
 - ⇒ At most n^2 checks
 - ⇒ Using BVH can save some checks

Types of Bounding Volumes

Many types of bounding volumes are possible





Figure 5.9: Four different kinds of bounding regions: (a) sphere, (b) axis-aligned bounding box (AABB), (c) oriented bounding box (OBB), and (d) convex hull. Each usually provides a tighter approximation than the previous one but is more expensive to test for overlapping pairs.


Each has benefits and issues

- ⇒ Spheres are simple and **orientation invariant** but do not fit tightly
- ⇒ AABBs are even simpler, but not orientation invariant, not tight
- ⇒ OBBs are orientation invariant, reasonably tight
- ⇒ Convex hulls are tight and orientation invariant, but require more computation


Probabilistic Roadmap in More Detail




Generating Random Samples




Rejecting Samples Outside \mathcal{C}_{free}




Collecting Enough Samples in C_{free}




Connect to k Nearest Neighbors (If Possible)




Connect to k Nearest Neighbors (If Possible)



Query Phase



Path Smoothing



PRM Algorithm – Roadmap Construction

First proposed by Kavraki et al.

Algorithm 6 Roadmap Construction Algorithm

Input:

n : number of nodes to put in the roadmap

k : number of closest neighbors to examine for each configuration

Output:

A roadmap $G = (V, E)$

```
1:  $V \leftarrow \emptyset$ 
2:  $E \leftarrow \emptyset$ 
3: while  $|V| < n$  do
4:   repeat
5:      $q \leftarrow$  a random configuration in  $\mathcal{Q}$ 
6:   until  $q$  is collision-free
7:    $V \leftarrow V \cup \{q\}$ 
8: end while
9: for all  $q \in V$  do
10:   $N_q \leftarrow$  the  $k$  closest neighbors of  $q$  chosen from  $V$  according to  $dist$ 
11:  for all  $q' \in N_q$  do
12:    if  $(q, q') \notin E$  and  $\Delta(q, q') \neq \text{NIL}$  then
13:       $E \leftarrow E \cup \{(q, q')\}$ 
14:    end if
15:  end for
16: end for
```

PRM Algorithm – Query Solving

Algorithm 7 Solve Query Algorithm

Input:

q_{init} : the initial configuration

q_{goal} : the goal configuration

k : the number of closest neighbors to examine for each configuration

$G = (V, E)$: the roadmap computed by algorithm 6

Output:


A path from q_{init} to q_{goal} or failure

| | |
|---|--|
| <pre>1: $N_{q_{\text{init}}} \leftarrow$ the k closest neighbors of q_{init} from V according to dist 2: $N_{q_{\text{goal}}} \leftarrow$ the k closest neighbors of q_{goal} from V according to dist 3: $V \leftarrow \{q_{\text{init}}\} \cup \{q_{\text{goal}}\} \cup V$ 4: set q' to be the closest neighbor of q_{init} in $N_{q_{\text{init}}}$ 5: repeat 6: if $\Delta(q_{\text{init}}, q') \neq \text{NIL}$ then 7: $E \leftarrow (q_{\text{init}}, q') \cup E$ 8: else 9: set q' to be the next closest neighbor of q_{init} in $N_{q_{\text{init}}}$ 10: end if 11: until a connection was succesful or the set $N_{q_{\text{init}}}$ is empty 12: set q' to be the closest neighbor of q_{goal} in $N_{q_{\text{goal}}}$</pre> | <pre>13: repeat 14: if $\Delta(q_{\text{goal}}, q') \neq \text{NIL}$ then 15: $E \leftarrow (q_{\text{goal}}, q') \cup E$ 16: else 17: set q' to be the next closest neighbor of q_{goal} in $N_{q_{\text{goal}}}$ 18: end if 19: until a connection was succesful or the set $N_{q_{\text{goal}}}$ is empty 20: $P \leftarrow$ shortest path($q_{\text{init}}, q_{\text{goal}}, G$) 21: if P is not empty then 22: return P 23: else 24: return failure 25: end if</pre> |
|---|--|

A Look at Completeness

Sampling-based algorithms are no longer **complete**!

- ⇒ If a solution exists, it will eventually find one and stop
- ⇒ When there is no solution, the algorithm may keep running forever (so we need to have a timeout for these methods)



We need a new notion of completeness

A New Notion of Completeness

Define a new notion of completeness based on **denseness** of sampling

- ⇒ A set of samples is **dense** if dispersion $\delta(P) \rightarrow 0$ as $|P| \rightarrow \infty$
- ⇒ This means that the roadmap will get into any opening
- ⇒ But it is hard to predict when if we do not know how big is the opening

Resolution completeness

- ⇒ For deterministic sampling (e.g., using a Halton sequence)
- ⇒ An algorithm is **resolution complete** if it samples deterministically and densely

Probabilistic completeness

- ⇒ For probabilistic methods
- ⇒ An algorithm is **probabilistic complete** if it samples probabilistically, e.g., uniformly random, and densely

