

CS 460/560

Introduction to Computational Robotics
Fall 2019, Rutgers University

Lecture 20

Aspects of Control

| | | | |
|----|----|----|----|
| 1 | 2 | 3 | 4 |
| 5 | 6 | 7 | 8 |
| 9 | 10 | 11 | 12 |
| 13 | 14 | 15 | 16 |

Instructor: Jingjin Yu

Outline

Feedback (closed-loop) control

- ⇒ Mathematical models of dynamical systems
- ⇒ Concept: open-loop v.s. closed-loop control
- ⇒ Ubiquity of feedback control systems
- ⇒ History of modern feedback control system: Watt's flyball governor
- ⇒ PID control
 - ⇒ Basic concepts
 - ⇒ Behavior of individual terms
 - ⇒ Tuning
- ⇒ Pure pursuit for controlling differential drive robots (DDR)

⇒ Optimal control

- ⇒ The Hamilton-Jacobi-Bellman equation, dynamic programming
- ⇒ The maximum principle
- ⇒ Time optimal trajectory of Dubin's car and DDR, bang-bang control

Modeling Dynamical Systems

A **dynamical system** (e.g., a car) is often modeled as

$$\dot{x} = f(x, u)$$

⇒ x : the **state** of the system, $x(t)$ yields the **trajectory**

⇒ E.g., for a car, $x(t) = (x_1(t), x_2(t), \theta(t))$

⇒ $\dot{x} = \frac{dx(t)}{dt}$ is the time derivative, i.e., the velocity of the system

⇒ For a car, $\dot{x} = (\dot{x}_1, \dot{x}_2, \dot{\theta})$

⇒ u : the **control input**

⇒ E.g., for a real car, approximately, $u = (\theta, v)$

⇒ θ is the front wheel bearing

⇒ v is the forward speed (for a 2-wheel drive, assuming no slippage)

⇒ u may be speed, acceleration, and so on...


⇒ f : system **evolution** function

⇒ How do x, u determine \dot{x}

In **discrete** settings, often written as $x_t = f(x_{t-1}, u_{t-1})$

⇒ May view this as integration of the continuous model: $x_t = x_{t-1} + \int_{t-1}^t \dot{x} dt$

⇒ Often written as $x_k = f(x_{k-1}, u_{k-1})$



Modeling Dynamical Systems, Continued

Examples

⇒ A car going at fixed speed along x_1 -axis: $\dot{x}_1 = 1$

⇒ In this case, $f(x, u) = 1$ is a constant

⇒ An accelerating car along x_1 -axis with acceleration a : $\dot{x}_1 = at$

⇒ $u = a$, the acceleration, $f(x, u) = at$, does not depend on x


⇒ A car going clockwise along the unit circle around the origin at unit speed

⇒ $\dot{x} = (\dot{x}_1, \dot{x}_2, \dot{\theta}) = (x_2, -x_1, -1)$

⇒ Initial condition: $x_1 = 1, x_2 = 0$

⇒ The car will keep circling the unit circle at unit speed

⇒ So it takes 2π time to go one round



Linear and non-linear systems

⇒ Linear systems: f is a linear function, e.g., $\dot{x} = Ax + Bu$

⇒ Non-linear systems: f is non-linear

What to grasp from the last two slides?

⇒ Dynamical systems may be modeled as we have described

⇒ In particular, given x_{k-1} , u_{k-1} , and $f(x, u)$, we can **predict** x_k

Open-Loop versus Closed-Loop Control

Open-loop control: the control input does not consider the current state of the system

⇒ Roughly speaking, this is saying u is independent of x in $\dot{x} = f(x, u)$

⇒ Example: $u = \dot{x} = 0$

⇒ Example: $u = (\dot{x}_1, \dot{x}_2) = (1, 1)$

⇒ For the car we work with, give constant input signals to the two wheels

Feedback (closed-loop) control: the control input takes into account the (observed) system state \tilde{x}

⇒ Example: damping $u = \dot{x} = -\tilde{x}$

⇒ What does this system do in one dimension?

⇒ If \tilde{x} is positive, \dot{x} is then negative, causing x to decrease. So it takes the system to the origin

⇒ Example: regulator $u = \dot{x} = c - (\tilde{x} - ct)$

⇒ What does this system do in one dimension?

⇒ If $\tilde{x} - ct$ is positive, meaning we have gone too fast, the system will slow down

⇒ Otherwise, the system will speed up

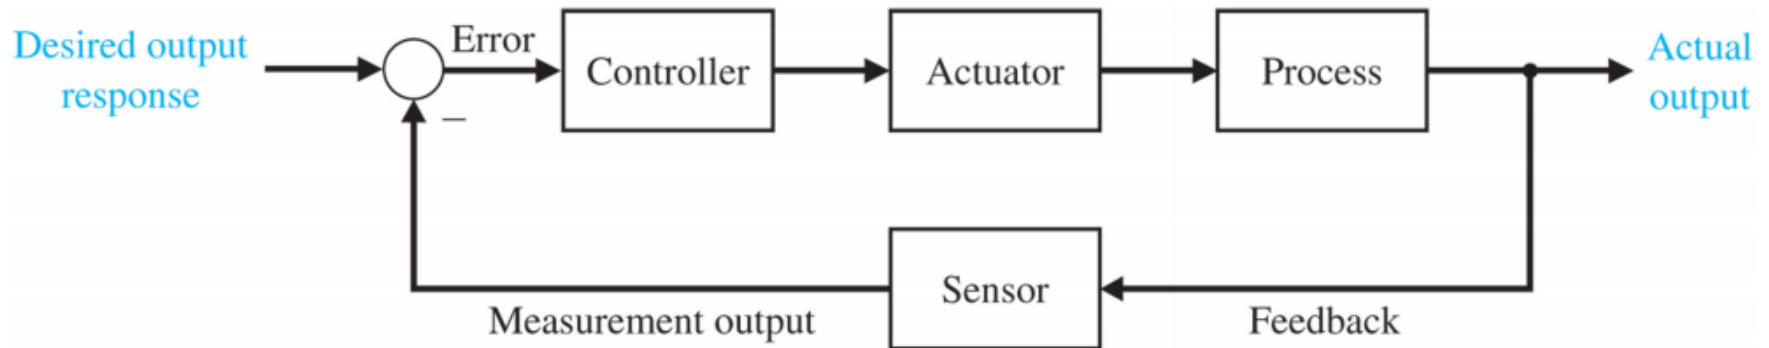
⇒ Overall, system goes at a speed of c

System Block Diagram

Open-loop control



Feedback (closed-loop) control



Real World Examples

Is a cannon open-loop or feedback-based?

⇒ Open-loop: we do not maintain control after shooting the shell

What about a missile?

⇒ Closed-loop: it tracks a target and is a form of PID control

In general, open-loop systems are cheaper than feedback-based system



Ubiquity of Feedback Control Systems

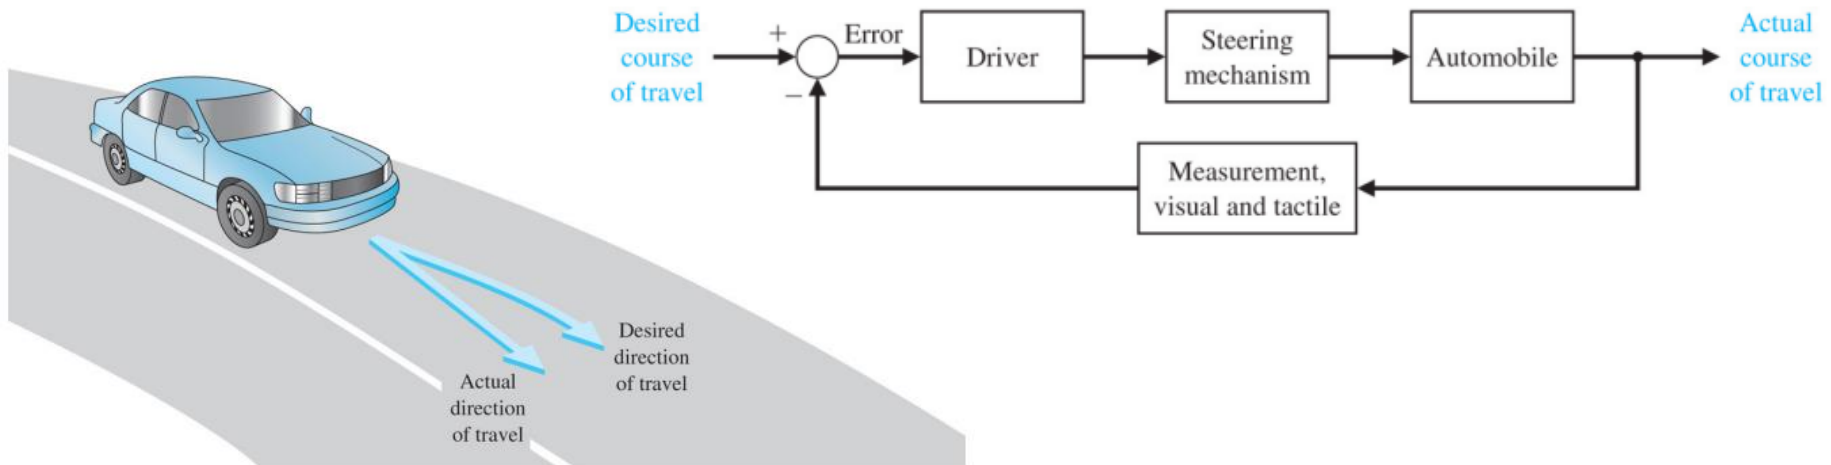
Feedback control systems are everywhere!

Example: ourselves

- ⇒ Walking with eyes open or closed
- ⇒ Reaction to electrical shock
- ⇒ Regulation of glucose (blood sugar) level with insulin (pancreas) and glycogen
 - ⇒ Sugar crash happens if you overload the system

Example: room temperature control (thermostat)


Example: driving a car



A Little History on Modern Feedback Control

After steam engine was invented, how to control its running speed is a problem of major interest

A successful design was Watt's flyball governor



PID Controller

PID controller stands for **proportional-integral-derivative controller**


⇒ There are many different “theoretical” feedback controllers

⇒ However, the final implementation often uses some form of PID control

General form:

$$u(t) = K_p e(t) + K_i \int_0^t e(\tau) d\tau + K_d \frac{de(t)}{dt}$$

Block diagram



PID Controller Breakdown

The error term $e(t)$

⇒ $e(t) = \text{Set Point} - \text{Current Location}$

⇒ E.g., if we want to go to the origin, may set $e(t) = (0,0) - (x_1(t), x_2(t))$

The proportional term, $K_p e(t)$

⇒ Adjust control based on **instant position error**

⇒ K_p : proportional gain, usually a positive constant

⇒ $e(t)$ large, then $u(t)$ is large

The integral term, $K_i \int_0^t e(\tau) d\tau$

⇒ Adjust system behavior based on **cumulative error**, slow response

⇒ Accelerates convergence to set point

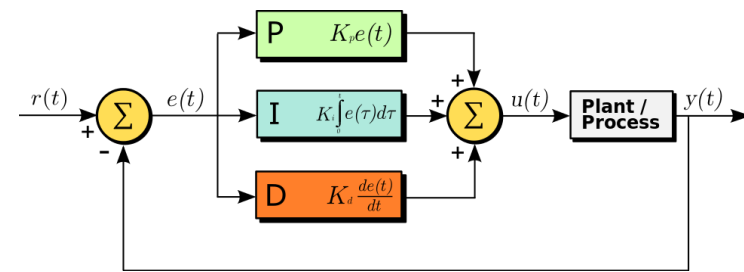
⇒ K_i : integral gain

The derivative term, $K_d \frac{de(t)}{dt}$

⇒ Adjust system behavior based on **differential error**, fast response (predictive)

⇒ Generally provides damping (preventing overshoot), improves stability


⇒ K_d : derivative gain




$$u(t) = K_p e(t) + K_i \int_0^t e(\tau) d\tau + K_d \frac{de(t)}{dt}$$

PID Controller Breakdown Example


Effects of varying the gains in a 1D system



Varying K_i



Varying K_p



Varying K_d


PID Controller Tuning

PID control applies easily to many systems since $e(t)$ can often be computed easily

However, the process of tuning a PID controller can be tricky



General (manual) method

- ⇒ Set $K_i = K_d = 0$ and tune K_p until the output oscillates
- ⇒ Adjust K_i so that the system converges to set point
- ⇒ Adjust K_d to remove oscillation until acceptable




Controlling a Quadcopter


How does a quadcopter work?



Up/down



Yaw change



Pitch/roll change

Quadcopter control

⇒ Hover

⇒ Basic trajectory following

⇒ Doing a sequence of “hovering” with new set points


Pure Pursuit for Differential Drive Robots

Most two wheeled robots can be viewed as a **differentially driven robot (DDR)**

⇒ Two wheel inputs in the range of $[-1, 1]$

Pure pursuit path following algorithm

- ⇒ From the current location of car, locate a waypoint of distance ℓ (some constant) on the desired trajectory
- ⇒ Compute the required curvature to the waypoint
- ⇒ Adjust wheel speeds to follow the computed arc
- ⇒ Note: the car's direction is tangential to the computed arc



Value Function and Principle of Optimality

Cost of trajectories can be captured with the **functional**

$$J(t, x, u) = \int_t^{t_1} L(s, x(s), u(s)) ds + K(x(t_1))$$

The **value function** is defined as the infimum of J over control u

$$V(t, x) := \inf_{u_{[t, t_1]}} J(t, x, u)$$


Principle of optimality states: For every $(t, x) \in [t_0, t_1) \times \mathbb{R}^n$ and every $\Delta t \in (0, t_1 - t]$, the value function satisfies

$$V(t, x) = \inf_{u_{[t, t+\Delta t]}} \left\{ \int_t^{t+\Delta t} L(s, x(s), u(s)) ds + V(t + \Delta t, x(t + \Delta t)) \right\}$$

Intuitive (provable)

But important

Yes, related to dynamic programming in CS



The Hamilton-Jacobi-Bellman (HJB) equation

From the principle of optimality, HJB equation can be derived

$$-\widehat{V}_t(t, x) = \inf_{u \in U} \{L(t, x, u) + \langle \widehat{V}_x(t, x), f(t, x, u) \rangle\}$$

- ⇒ A PDE providing **necessary** and **sufficient** conditions for optimal control u^*
- ⇒ From u^* one can also obtain the optimal trajectory x^*
- ⇒ From HJB, can derive conditions for optimal solutions

E.g., a simple integrator $\dot{x} = u$ with $L(x, u) = x^2 + u^3$. HJB

$$-V_t(t, x) = \inf_{u \in \mathbb{R}} \{x^2 + u^3 + V_x(t, x)u\}$$

Optimal control:

$$u^*(t) = -\sqrt{\frac{1}{3} V_x(t, x)}$$

PDE:

$$-V_t(t, x) = x^2 - 2 \left(\frac{1}{3} V_x(t, x) \right)^{\frac{3}{2}}$$

In general, closed form solutions for HJB are difficult to come by

Pontryagin's Maximum Principle

HJB provides necessary and sufficient conditions for optimality

⇒ But does so requiring the value function $V(t, x)$ be C^1

⇒ That is, first order partial derivatives must be continuous

Pontryagin's maximum principle address this.

⇒ The (fixed end point problem) setup: $J(u) = \int_{t_0}^{t_f} L(x, u) dt + K(t_f, x_f)$ and $\dot{x} = f(x, u)$

⇒ Hamiltonian: $H(x, u, p, p_0) := \langle p, f(x, u) \rangle + p_0 L(x, u)$

⇒ Maximum principle says that

⇒ There exist $(p^*, p_0) \neq (0, 0)$ such that $\dot{x}^* = H_p(x^*, u^*, p^*, p_0^*)$, $\dot{p}^* = -H_x(x^*, u^*, p^*, p_0^*)$

⇒ There exists global maximum $H(x^*(t), u^*(t), p^*(t), p_0^*) \geq H(x^*(t), u(t), p^*(t), p_0^*)$

⇒ $H(x^*(t), u^*(t), p^*(t), p_0^*) = 0$ for all $t \in [t_0, t_f]$

A bit of history

⇒ HJB equation: 1957, from US

⇒ Maximum principle: 1956, from USSR

⇒ The cold war era...

Optimal Trajectories for Dubins Car, DDR

Bang-bang control

- ⇒ Both HJB and the maximum principle lead to **bang-bang** control
- ⇒ In general, optimal trajectory uses extreme control inputs

E.g., moving from $x = 0$ to $x = 1$ with $\dot{x} = u \in [-1, 1]$

- ⇒ What is the time optimal strategy?
- ⇒ Move with $\dot{x} = 1$
- ⇒ What if $\ddot{x} = u \in [-1, 1]$?
- ⇒ Move with $\ddot{x} = 1$ halfway, then $\ddot{x} = -1$

E.g. Dubins car

- ⇒ Three types of moves: L, S, R
- ⇒ A distance optimal solution uses at most a sequence of three
- ⇒ Similar results applies to DDR, e.g. extremal

